

Black-box Adversarial Attacks on Video Recognition Models

Linxi Jiang^{1*}, Xingjun Ma^{2*}, Shaoxiang Chen¹, James Bailey², Yu-Gang Jiang^{1**}

¹Shanghai Key Lab of Intelligent Information Processing, School of Computer Science, Fudan University

² School of Computing and Information Systems, The University of Melbourne

{lxjiang18, sxchen13, ygj}@fudan.edu.cn, {xingjun.ma, baileyj}@unimelb.edu.au

ABSTRACT

Deep neural networks (DNNs) are known for their vulnerability to adversarial examples. These are examples that have undergone small, carefully crafted perturbations, and which can easily fool a DNN into making misclassifications at test time. Thus far, the field of adversarial research has mainly focused on image models, under either a white-box setting, where an adversary has full access to model parameters, or a black-box setting where an adversary can only query the target model for probabilities or labels. Whilst several white-box attacks have been proposed for video models, black-box video attacks are still unexplored. To close this gap, we propose the first black-box video attack framework, called V-BAD. V-BAD utilizes *tentative perturbations* transferred from image models and *partition-based rectifications* found by the NES to obtain good adversarial gradient estimates with fewer queries to the target model. V-BAD is equivalent to estimating the projection of the adversarial gradient on a selected subspace. Using three benchmark video datasets, we demonstrate that V-BAD can craft both untargeted and targeted attacks to fool two state-of-the-art deep video recognition models. For the targeted attack, it achieves $>93\%$ success rate using only an average of $3.4 \sim 8.4 \times 10^4$ queries, a similar number of queries to state-of-the-art black-box image attacks. This is despite the fact that videos often have two orders of magnitude higher dimensionality than static images. We believe that V-BAD is a promising new tool to evaluate and improve the robustness of video recognition models to black-box adversarial attacks.

CCS CONCEPTS

- Computing methodologies → Activity recognition and understanding.

KEYWORDS

Adversarial examples; video recognition; black-box attack; model security

ACM Reference Format:

Linxi Jiang^{1*}, Xingjun Ma^{2*}, Shaoxiang Chen¹, James Bailey², Yu-Gang Jiang^{1**}. 2019. Black-box Adversarial Attacks on Video Recognition Models. In *Proceedings of the 27th ACM International Conference on Multimedia (MM '19)*, October 21–25, 2019, Nice, France. ACM, New York, NY, USA, 9 pages. <https://doi.org/10.1145/3343031.3351088>

* indicates equal contributions. ** indicates corresponding author.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

MM '19, October 21–25, 2019, Nice, France

© 2019 Association for Computing Machinery.

ACM ISBN 978-1-4503-6889-6/19/10...\$15.00

<https://doi.org/10.1145/3343031.3351088>



Figure 1: An example of black-box video adversarial attacks (targeted). The original video (top) can be correctly recognized while the adversarial one (bottom) generated by our proposed method is misclassified by the same video model.

'19, October 21–25, 2019, Nice, France. ACM, New York, NY, USA, 9 pages. <https://doi.org/10.1145/3343031.3351088>

1 INTRODUCTION

Deep Neural Networks (DNNs) are a family of powerful models that have demonstrated superior performance in a wide range of visual understanding tasks that have been extensively studied in both the multimedia and computer vision communities such as video recognition [13–15, 17, 42], image classification [20, 45] and video captioning [25, 43]. Despite their current success, DNNs have been found to be extremely vulnerable to adversarial examples (or attacks) [9, 37]. For DNN classifiers, adversarial examples can be easily generated by applying adversarial perturbations to clean (normal) examples towards maximizing the classification error [3, 9, 30]. For images, the perturbations are often small and visually imperceptible to human observers, but they can fool DNNs into making misclassifications with high confidence. The vulnerability of DNNs to adversarial examples has raised serious security concerns for their deployment in security-critical applications, such as face recognition [22] and self-driving cars [8]. Hence, the study of adversarial examples for DNNs has become a crucial task for secure deep learning.

Adversarial examples can be generated by an attack method (also called an adversary) following either a white-box setting (white-box attacks) or a black-box setting (black-box attacks). In the white-box setting, an adversary has full access to the target model (the model to attack), including model parameters and training settings. In the black-box setting, an adversary only has partial information about the target model, such as the labels or probabilities output by the model. White-box methods generate an adversarial example by

applying one step or multiple steps perturbations on a clean test sample, following the direction of the adversarial gradient [3, 9]. The adversarial gradient is the gradient of an adversarial loss, which is typically defined to maximize (rather than minimize) classification error. However, in the black-box setting, adversarial gradients are not accessible to an adversary. In this case, the adversary can first attack a local surrogate model and then transfer these attacks to the target model [28, 31, 32]. Alternatively they may use a black-box optimization method such as Finite Differences (FD) or Natural Evolution Strategies (NES), to estimate the gradient [2, 6, 12].

A number of attack methods have been proposed [3, 27, 30], however, most of them focus on either image models, or video models but in a white-box setting [24, 34]. Different from these works, in this paper, we propose a framework for the generation of adversarial attacks against video recognition models, specifically in a black-box setting. Significant progress has been achieved for black-box image attacks, but not for black-box video attacks. A key reason is that videos typically have much higher dimensionality (often two magnitudes higher) than static images. On static images, for most attacks to succeed, existing black-box methods must use $\sim 10^4$ queries [2, 6] on CIFAR-10 [19] images, and $\sim 10^5$ queries [12] on ImageNet [7] images. Due to their massive input dimensions, black-box attacks on videos generally require two orders of magnitude more queries for gradient estimation than are needed for images. This makes black-box video attacks impractical, taking into account time and budget constraints. To better evaluate the robustness of video models, it is therefore important to explore efficient black-box methods that can generate attacks using fewer queries.

In this paper, we propose a simple and efficient framework for the generation of black-box adversarial attacks on video recognition models. Intuitively, we exploit the transferability of adversarial perturbations and the derivative-free optimization methods to obtain accurate estimation of the adversarial gradient. In particular, we first generate tentative perturbations as a rough estimate of the true adversarial gradient using ImageNet-pretrained DNNs. We then rectify these tentative perturbations in *patches* (or partitions) using NES, by querying the target model. Our proposed framework only needs to estimate a small number of directional derivatives (of the patches) rather than estimating pixel-wise derivatives, making it an efficient framework for black-box video attacks. Figure 1 shows an example of video adversarial attacks generated by our proposed method. In summary, our main contributions are:

- We study the problem of black-box attacks on video recognition models and propose a general framework called V-BAD, to generate black-box video adversarial examples. To the best of our knowledge, our proposed framework V-BAD is the first black-box adversarial attack framework for videos.
- Our proposed framework V-BAD exploits both the transferability of adversarial perturbations via the use of tentative perturbations, and the advantage of NES gradient estimation for patch-level gradient rectification. We also show that V-BAD is equivalent to estimating the projection of the adversarial gradient on a selected subspace.
- We conduct an empirical evaluation using three benchmark video datasets and two state-of-the-art video recognition models. We show that V-BAD can achieve high attack success

rates with few queries to target models, making it a useful tool for the robustness evaluation of video models.

2 RELATED WORK

White-box Image Attack. The fast gradient sign method (FGSM) [9] crafts an adversarial example by perturbing a normal sample along the gradient direction towards maximizing the classification error. FGSM is a single-step attack, and can be applied iteratively to improve adversarial strength [22]. Projected Gradient Descent (PGD) [27] is another iterative method that is regarded as the strongest first-order attack. PGD projects the adversarial example back onto the ϵ -ball centered at the original sample when perturbations go beyond the ϵ -ball. The C&W attack solves the attack problem via an optimization framework [3], and is arguably the state-of-the-art white-box attack. There also exists other types of white-box methods, e.g., Jacobian-based Saliency Map Attack (JSMA) [33], DeepFool [29] and elastic-net attack (EAD) [5].

Black-box Image Attack. In the black-box setting, the adversarial gradients are not directly accessible by an adversary. As such, black-box image attacks either exploit the transferability of adversarial examples (also known as transferred attacks) or make use of gradient estimation techniques. It was first observed in [37] that adversarial examples are transferable across models, even if the models have different architectures or were trained separately. Papernot et al. [32] trained a surrogate model locally on synthesized data, with labels obtained by querying the target model. It then generates adversarial examples from the surrogate model using white-box methods to attack the target model. However, training a surrogate model on synthesized data often incurs a huge number of queries, and the transferability of generated adversarial examples is often limited. Chen et al. [6] proposed the use of Finite Differences (FD), a black-box gradient estimation method, to estimate the adversarial gradient. Bhagoji et al. [2] accelerated FD-based gradient estimation with dimensionality reduction techniques such as PCA. Compared with FD, Ilyas et al. [12] demonstrated improved performance with fewer queries by using the other type of gradient estimation method called Natural Evolutionary Strategies (NES).

White-box Video Attack. In contrast to image adversarial examples, much less work has been done for video adversarial examples. White-box video attacks were first investigated in [40], which discussed the sparsity and propagation of adversarial perturbations across video frames. Li et al. [24] leveraged Generative Adversarial Networks (GANs) to perturb each frame in real-time video classification. In this paper, we explore black-box attacking methods against state-of-the-art video recognition models, which, to the best of our knowledge, is the first work on black-box video attacks.

Video Recognition Models. CNN+LSTM based models [16, 44] exploited the temporal information contained in successive frames, with recurrent layers capturing long term dependencies on top of CNNs. Carreira and Zisserman [4] proposed the Inflated 3D ConvNet(I3D) with Inflated 2D filters and pooling kernels of traditional 2D CNNs and learn hierarchical spatio-temporal representations directly from videos. In this paper, we use two representative state-of-the-art video recognition models, CNN+LSTM and I3D, as our target models to attack.

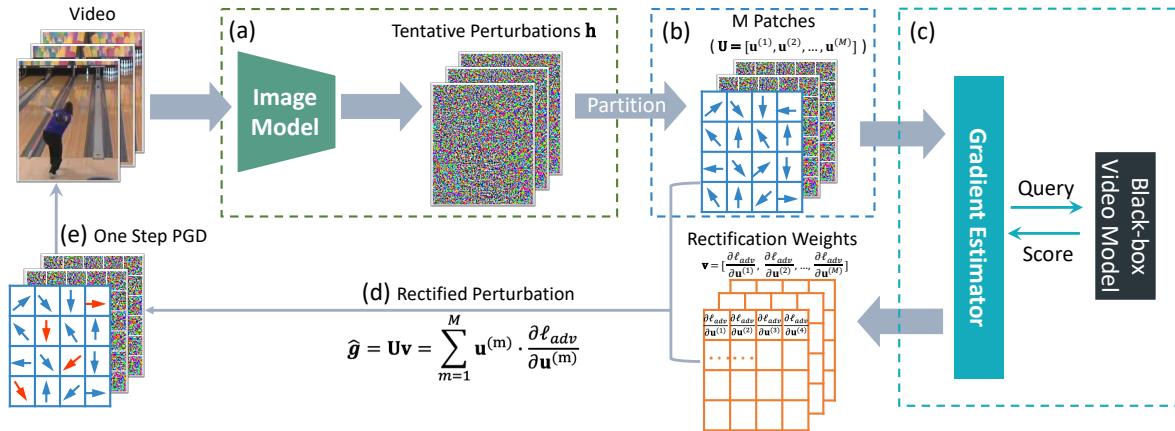


Figure 2: Overview of the proposed V-BAD framework for black-box video attacks. At step (a), V-BAD passes video frames into a public image model to obtain tentative perturbations. At steps (b)-(d), V-BAD partitions tentative perturbations into patches and perform patch-wise rectification. At step (e), a one-step PGD is performed according to the rectified perturbations. Steps (a)-(d) are proposed to accelerate the attack. The process will be performed iteratively until an adversarial example is found.

3 PROPOSED FRAMEWORK V-BAD

3.1 Preliminaries

We denote a video sample by $\mathbf{x} \in \mathcal{X} \subset \mathbb{R}^{N \times H \times W \times C}$ with N , H , W , C denoting the number of frames, frame height, frame width, and the number of channels respectively, and its associated true class by $y \in \mathcal{Y} = \{1, \dots, K\}$. Video recognition is to learn a DNN classifier $f(\mathbf{x}; \theta) : \mathcal{X} \rightarrow \mathcal{Y}$ by minimizing the classification loss $\ell(f(\mathbf{x}; \theta), y)$, and θ denotes the parameters of the network. When the context is clear, we abbreviate $\ell(f(\mathbf{x}; \theta), y)$ as $\ell(\mathbf{x}, y)$. The goal of adversarial attack is to find an adversarial example \mathbf{x}_{adv} that can fool the network to make a false prediction, while keeping the adversarial example \mathbf{x}_{adv} within the ϵ_{adv} -ball centered at the original example \mathbf{x} ($\|\mathbf{x}_{adv} - \mathbf{x}\|_p \leq \epsilon_{adv}$). Following early works [1, 12, 26, 27, 39], in this paper, we only focus on the L_∞ -norm, that is, $\|\mathbf{x}_{adv} - \mathbf{x}\|_\infty \leq \epsilon_{adv}$.

There are two types of adversarial attacks: untargeted attack and targeted attack. Untargeted attack is to find an adversarial example that can be misclassified as an arbitrary class other than the correct one (e.g. $f(\mathbf{x}_{adv}) \neq y$), while targeted attack is to find an adversarial example that can be misclassified as a targeted adversarial class (e.g. $f(\mathbf{x}_{adv}) = y_{adv}$ and $y_{adv} \neq y$). For simplicity, we denote the adversarial loss function that should be optimized to find an adversarial example by $\ell_{adv}(\mathbf{x})$, and let $\ell_{adv}(\mathbf{x}) = -\ell(\mathbf{x}, y)$ for untargeted attack and $\ell_{adv}(\mathbf{x}) = \ell(\mathbf{x}, y_{adv})$ for targeted attack. We also denote the adversarial gradient of the adversarial loss to the input as $\mathbf{g} = \nabla_{\mathbf{x}} \ell_{adv}(\mathbf{x})$. Accordingly, an attacking method is to minimize the adversarial loss $\ell_{adv}(\mathbf{x})$ by iteratively perturbing the input sample following the direction of the adversarial gradient \mathbf{g} .

Threat Model. Our threat model follows the query-limited black-box setting as follows. The adversary takes the video classifier f as a black-box oracle and only has access to its output of the top 1 score. More specifically, during the attacking process, given a clean sample \mathbf{x} , the adversary can query the target model f to obtain the top 1 label \hat{y} and its probability $P(\hat{y}|\mathbf{x})$. The adversary is asked to generate attacks within Q number of queries. We consider both untargeted and targeted black-box attacks.

3.2 Framework Overview

The structure of the proposed framework, namely V-BAD, for black-box video attacks is illustrated in Figure 2. Following steps (a)-(e) highlighted in the figure, V-BAD perturbs an input video iteratively as follow: (a) It passes video frames into a public image model (such as ImageNet-pretrained networks) to obtain pixel-wise *tentative perturbations* \mathbf{h} ; (b) It then partitions \mathbf{h} into a set of M patches $\mathbf{U} = [\mathbf{u}^{(1)}, \mathbf{u}^{(2)}, \dots, \mathbf{u}^{(M)}]$ where $\mathbf{u}^{(m)}$ represents the m -th patch; (c) A black-box gradient estimator estimates the rectification weight $\mathbf{v}_m = \frac{\partial \ell_{adv}}{\partial \mathbf{u}^{(m)}}$ required for each patch, via querying the target video model; (d) The patch-wise rectification weights are applied on the patches to obtain the rectified pixel-wise perturbations $\hat{\mathbf{g}}$; (e) Apply one step PGD perturbation on the input video according to $\hat{\mathbf{g}}$. Specifically, the proposed method at the t -th PGD perturbation step can be described as:

$$\mathbf{v} = \mathbf{0} \quad (1)$$

$$\mathbf{h}^t = \phi(\mathbf{x}^{t-1}) \quad (2)$$

$$\mathbf{U} = G(\mathbf{h}^t) \quad (3)$$

$$\hat{\mathbf{v}} = \mathbf{v} + \nabla_{\mathbf{v}} \ell_{adv}(\mathbf{x}^{t-1} + R(\mathbf{v}, \mathbf{U})) \quad (4)$$

$$\hat{\mathbf{g}} = R(\hat{\mathbf{v}}, \mathbf{U}) \quad (5)$$

$$\mathbf{x}^t = \Pi_\epsilon(\mathbf{x}^{t-1} - \alpha \cdot \text{sign}(\hat{\mathbf{g}})) \quad (6)$$

where, \mathbf{x}^{t-1} is the adversarial example generated at step $t-1$, $\phi(\cdot)$ is the function to extract tentative perturbations, $G(\cdot)$ is a partitioning method that splits pixels of \mathbf{h}^t into a set of patches \mathbf{U} , $\hat{\mathbf{v}}$ is the estimated rectification weights for patches in \mathbf{U} via a black-box gradient estimator, $R(\mathbf{v}, \mathbf{U})$ is a rectification function that applies patch-wise rectifications $\hat{\mathbf{v}}$ to patches in \mathbf{U} to obtain pixel-wise rectified perturbations $\hat{\mathbf{g}}$, $\Pi(\cdot)$ is a projection operation, $\text{sign}(\cdot)$ is the sign function, α is the PGD step size, ϵ is the perturbation constraint. $R(\mathbf{v}, \mathbf{U})$ applies the rectification weights to patches, thus enables us to estimate the gradients with respect to the patches in place of the gradients with respect to the raw pixels, which reduces the dimension of attack space from $\mathbb{R}^{N \times H \times W \times C}$ to \mathbb{R}^M where M

Algorithm 1 Targeted V-BAD attack

Input: Top-1 probability $P(\hat{y}|\mathbf{x})$ with respect to classifier f , target class y_{adv} and video \mathbf{x}
Output: Adversarial video \mathbf{x}_{adv} with $\|\mathbf{x}_{adv} - \mathbf{x}\|_\infty \leq \epsilon_{adv}$
Parameters: Perturbation bound ϵ_{adv} , epsilon decay $\Delta\epsilon$, PGD step size α

$$\epsilon \leftarrow 1, \mathbf{x}_{adv} \leftarrow \text{video of target class } y_{adv}$$

while $\epsilon > \epsilon_{adv}$ **do**

- $\hat{\mathbf{g}} \leftarrow \text{Gradient estimation following Eq. (1)-(5)}$
- $\hat{\epsilon} \leftarrow \epsilon - \Delta\epsilon$
- $\hat{\mathbf{x}}_{adv} \leftarrow \text{CLIP}(\mathbf{x}_{adv} - \alpha \cdot \hat{\mathbf{g}}, \mathbf{x} - \hat{\epsilon}, \mathbf{x} + \hat{\epsilon})$
- if** $y_{adv} = \text{Top-1}(P(\cdot|\hat{\mathbf{x}}_{adv}))$ **then**

 - $\mathbf{x}_{adv} \leftarrow \hat{\mathbf{x}}_{adv}, \epsilon \leftarrow \hat{\epsilon}$

- else**

 - $\hat{\mathbf{x}}_{adv} \leftarrow \text{CLIP}(\mathbf{x}_{adv} - \alpha \cdot \hat{\mathbf{g}}, \mathbf{x} - \epsilon, \mathbf{x} + \epsilon)$
 - if** $y_{adv} = \text{Top-1}(P(\cdot|\hat{\mathbf{x}}_{adv}))$ **then**

 - $\mathbf{x}_{adv} \leftarrow \hat{\mathbf{x}}_{adv}$

- end if**

end while

return \mathbf{x}_{adv}

is the number of patches. Each of the above operations will be explained in the following sections.

Targeted V-BAD Attack. For a targeted attack, we need to ensure that the target class remains in the top-1 class during the attacking process, as the score of the target class is required for gradient estimation. Thus, instead of the original sample \mathbf{x} , we begin with a sample from the target class (e.g. $\mathbf{x}^0 = \mathbf{x}'$ and $f(\mathbf{x}') = y_{adv}$), then gradually (step by step) reduce the perturbation constraint ϵ from 1 (for normalized inputs $\mathbf{x} \in [0, 1]$) to ϵ_{adv} while maintaining the targeted class as the top-1 class. Note that although we begin with $\mathbf{x}^0 = \mathbf{x}'$, the adversarial example will eventually be moved into the ϵ_{adv} -ball centered at the original example \mathbf{x} . The targeted V-BAD attack is described in Algorithm 1. We use an epsilon decay $\Delta\epsilon$ to control the reduction size of the perturbation bound. Note that the epsilon decay $\Delta\epsilon$ and PGD step size α are dynamically adjusted and we omit the details for simplicity.

Untargeted V-BAD Attack. For untargeted attack, we use the original clean example as our starting point: $\mathbf{x}^0 = \mathbf{x}$, where \mathbf{x} is the original clean example. And the perturbation bound ϵ is set to the constant ϵ_{adv} throughout the attacking process. Given estimated adversarial gradient $\hat{\mathbf{g}}$, untargeted attack follows the standard procedure of PGD attack, therefore we omit the algorithm here.

3.3 Tentative Perturbations

Tentative perturbations refer to initialized pixel-wise adversarial perturbations generated at the beginning of each perturbation step, via function $\phi(\cdot)$ in Eq. (2). As PGD only needs the sign of adversarial gradient to perturb a sample, here we also take the sign values for tentative perturbations.

Here, we discuss three types of tentative perturbations. 1) Random: the perturbation at each input dimension takes value 1 or -1 equiprobably. Random perturbations only provide a stochastic exploration of the input space, which can be extremely inefficient

due to the massive input dimensions of videos. 2) Static: the perturbation for each input dimension is fixed to 1. Fixed perturbations impose a strong constraint on the exploration space, with the same tentative perturbations used for each gradient estimation. 3) Transferred: tentative perturbations can alternatively be transferred from off-the-shelf pre-trained image models. Since natural images share certain similar patterns, image adversarial examples often transfer across models or domains, though the transferability can be limited and will vary depending on the content of the images. To better exploit such transferability for videos, we propose to white-box attack a pre-trained image model such as an ImageNet pre-trained DNN to extract tentative perturbations for each frame of the video. The resulting transferred perturbations can provide useful guidance for space exploration and help reduce the number of queries.

The tentative perturbations for an intermediate perturbed sample \mathbf{x}^{t-1} can be extracted by:

$$\phi(\mathbf{x}^{t-1}) = \text{sign}\left(\nabla_{\mathbf{x}^{t-1}} \frac{1}{N} \sum_{n=1}^N \|\mathbf{b}_n \circ \tilde{f}_l(\mathbf{x}_n^{t-1}) - \mathbf{b}_n \circ \mathbf{r}_n\|_2^2\right), \quad (7)$$

where $\text{sign}(\cdot)$ is the sign function, \circ is the element-wise product, $\|\cdot\|_2^2$ is the squared L_2 -norm, N is the total number of frames, $\tilde{f}_l(\cdot)$ is the l -th layer output (e.g. deep features) of a public image model \tilde{f} , \mathbf{b}_n is a random mask on the n -th frame, and \mathbf{r}_n is a target feature map which has the same dimensions as $\tilde{f}_l(\mathbf{x}_n^{t-1})$. $\phi(\cdot)$ is to generate tentative perturbations by minimizing the L_2 distance between the feature map $\tilde{f}_l(\cdot)$ and a target (or adversarial) feature map \mathbf{r}_n . For untargeted attack, \mathbf{r}_n is a feature map of Gaussian random noise, while for targeted attack, \mathbf{r}_n is the feature map $\tilde{f}_l(\mathbf{x}_n)$ of the n -th frame of a video \mathbf{x}' from the target class y_{adv} . The random mask \mathbf{b}_n is a pixel-wise mask with each element randomly set to either 1 or 0. The use of random mask can provide some randomness to the exploration of the attack space, and help avoid getting stuck in situations where the frame-wise tentative perturbations are extremely inaccurate to allow any effective rectification.

3.4 Partition-based Rectification

Tentative perturbations are transferred rough estimate of the true adversarial gradient, thus it should be further rectified to better approximate the true adversarial gradient. To achieve this more efficiently, we propose to perform rectifications at the patch-level (in contrast to pixel-level) and use gradient estimation methods to estimate the proper rectification weight for each patch. That is, we first partition tentative perturbations into separate parts, which we call perturbation patches, and then estimate the rectification weights for these patches. In this subsection, we will introduce the partitioning methods, rectification function, and rectification weights estimation successively.

Partitioning Methods. We consider two types of partitioning strategies (e.g. function $G(\cdot)$ in Eq. (3)). 1) *Random*: dividing input dimensions randomly into a certain number of patches. Note that the input dimensions within one patch can be nonconsecutive, although we refer to it as a patch. Random partitioning does not consider the local correlations between input dimensions but can be used as a baseline to assess the effectiveness of carefully designed partitioning methods. 2) *Uniform*: splitting a frame uniformly into

Algorithm 2 NES Estimation of Patch Rectification

Input: Adversarial loss $\ell_{adv}(\mathbf{x})$ with respect to input \mathbf{x} and tentative perturbation patches \mathbf{U}
Output: Estimate of $\nabla_{\mathbf{v}} \ell_{adv}(\mathbf{x} + R(\mathbf{v}, \mathbf{U}))$
Parameters: Search variance σ , sampling size λ

```

 $\hat{\mathbf{v}} \leftarrow \mathbf{0}_M$ 
for  $k = 1$  to  $\lambda/2$  do
     $\delta_k \leftarrow \mathcal{N}(\mathbf{0}_M, I_{M \times M})$ 
     $\hat{\mathbf{v}} \leftarrow \hat{\mathbf{v}} + \delta_k \cdot \text{TransformedAdvLoss}(\ell_{adv}(\mathbf{x} + R(\mathbf{v} + \sigma \delta_k, \mathbf{U})))$ 
     $\hat{\mathbf{v}} \leftarrow \hat{\mathbf{v}} - \delta_k \cdot \text{TransformedAdvLoss}(\ell_{adv}(\mathbf{x} - R(\mathbf{v} + \sigma \delta_k, \mathbf{U})))$ 
end for
 $\hat{\mathbf{v}} \leftarrow \frac{1}{\lambda \sigma} \hat{\mathbf{v}}$ 
return  $\hat{\mathbf{v}}$ 

```

a certain number of patches. This will produce frame patches that preserve local dimensional correlations.

Each patch constitutes a vector with the same number of dimensions as the input by zero padding: assigning zero values to dimensions that do not belong to the patch, while keeping its values at dimensions that belong to the patch. That is, for the m -th patch, we have a vector $\mathbf{u}^{(m)}$:

$$\mathbf{u}_j^{(m)} = \begin{cases} \mathbf{h}_j & \text{if dimension } j \text{ is in the } m\text{-th patch} \\ 0 & \text{otherwise.} \end{cases}$$

Note that, $\mathbf{u}^{(m)}$ has equal number of dimensions to both \mathbf{x} and \mathbf{h} , namely, $\mathbf{u}^{(m)} \in \mathbb{R}^{N \times H \times W \times C}$. We normalize the $\mathbf{u}^{(m)}$ to unit vector $\mathbf{u}^{(m)} = \frac{\mathbf{u}^{(m)}}{\|\mathbf{u}^{(m)}\|}$ ($\|\cdot\|$ is the vector norm), which we call the direction vector of the m -th patch. The partition function can be written as:

$$G(\mathbf{h}) = [\mathbf{u}^{(1)}, \mathbf{u}^{(2)}, \dots, \mathbf{u}^{(M)}].$$

We denote the output of partition function as $\mathbf{U} = G(\mathbf{h})$.

Rectification Function. The rectification function $R(\mathbf{v}, \mathbf{U})$ (in Eq. (4) and (5)) can be expressed as applying the components of rectification vector \mathbf{v} as weights to each patch. Given a patch-wise rectification vector \mathbf{v} (estimated by some black-box gradient estimator), rectification is to apply patch weights to the direction vectors of the patches:

$$R(\mathbf{v}, \mathbf{U}) = \mathbf{U}\mathbf{v} = [\mathbf{u}^{(1)}, \mathbf{u}^{(2)}, \dots, \mathbf{u}^{(M)}] \cdot (\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_M)^T.$$

With the above partition and rectification functions, we can optimize the rectification weights \mathbf{v} over patches instead of all input dimensions. This effectively reduces the dimensionality of the exploration space from $N \times H \times W \times C$ to M .

Rectification Weights Estimation. For the estimation of rectification weights for patches, we propose to use the NES estimator which has been shown to be more efficient than FD in black-box image attacks [12]. We will discuss the efficiency of FD compared with NES in Section 4. Instead of maximizing the adversarial objective directly, NES maximizes the expected value of the objective under a search distribution. Different from [12] where the gradient of adversarial loss with respect to the input is estimated, we estimate the gradient with respect to the patch weights \mathbf{v} . For an adversarial loss function $\ell_{adv}(\cdot)$, input \mathbf{x} , tentative patches \mathbf{U} , current parameters

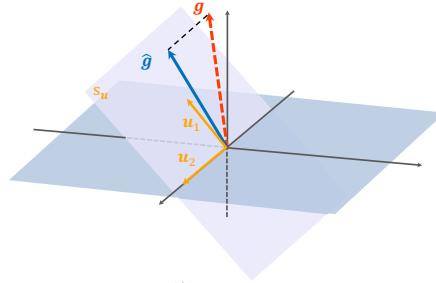


Figure 3: Projection ($\hat{\mathbf{g}}$) of gradient \mathbf{g} on subspace s_u .

\mathbf{v} and a search distribution $\pi(\mathbf{y}|\mathbf{v})$, we have:

$$\nabla_{\mathbf{v}} \mathbb{E}_{\pi(\mathbf{y}|\mathbf{v})} [\ell_{adv}(\mathbf{x} + R(\mathbf{y}, \mathbf{U}))] = \mathbb{E}_{\pi(\mathbf{y}|\mathbf{v})} [\ell_{adv}(\mathbf{x} + R(\mathbf{y}, \mathbf{U})) \nabla_{\mathbf{v}} \log(\pi(\mathbf{y}|\mathbf{v}))].$$

Following [12, 35, 41], we use the normal distribution as the search distribution, that is, $\mathbf{y} = \mathbf{v} + \sigma \boldsymbol{\delta}$ where σ is the search variance and $\boldsymbol{\delta} \sim \mathcal{N}(0, I)$ (standard normal distribution). We use antithetic sampling to generate a population of λ number of δ_k values: first sample Gaussian noise for $k \in \{1, \dots, \frac{\lambda}{2}\}$, then set $\delta_k = -\delta_{\lambda-k+1}$ for $k \in \{(\frac{\lambda}{2}+1), \dots, \lambda\}$. Evaluating the gradient with a population of λ points sampled under this scheme yields the following gradient estimate:

$$\nabla_{\mathbf{v}} \mathbb{E} [\ell_{adv}(\mathbf{x} + R(\mathbf{y}, \mathbf{U}))] \approx \frac{1}{\lambda \sigma} \sum_{k=1}^{\lambda} \delta_k \ell_{adv}(\mathbf{x} + R(\mathbf{v} + \sigma \delta_k, \mathbf{U})).$$

Within each perturbation step, \mathbf{v} is initialized to zero values (see Eq. (1)) such that the starting point $\mathbf{x} + R(\mathbf{v}, \mathbf{U})$ is centered at \mathbf{x} .

The complete estimation algorithm for patch rectification is described in Algorithm 2, where the function $\text{TransformedAdvLoss}(\cdot)$ is a ranking-based nonlinear transformation on the adversarial loss. The transformed loss increases monotonically with the original adversarial loss [41]. For targeted attacks, it also “punishes” δ_k that fails to maintain the target class as the top-1 class by the highest loss value. With this estimated gradient for \mathbf{v} , we can update the rectification weights as $\hat{\mathbf{v}} = \mathbf{v} + \nabla_{\mathbf{v}} \mathbb{E} [\ell_{adv}(\mathbf{x} + R(\mathbf{y}, \mathbf{U}))]$. We then use $\hat{\mathbf{v}}$ to rectify tentative perturbations \mathbf{U} to $R(\hat{\mathbf{v}}, \mathbf{U})$, which allows us to apply one step of PGD perturbation following Eq. (6).

3.5 Analysis of V-BAD

We prove that patch-rectified perturbations is equivalent to estimating the projected adversarial gradient on a selected subspace. Let \mathcal{S} be the input space with respect to X , direction vectors (e.g. \mathbf{u}) of the tentative perturbation patches define a vector subspace $s_u \subseteq \mathcal{S}$. Using the chain rule, we can show that the gradients with respect to the patch-wise rectification weights \mathbf{v} are the directional derivatives of the patch directions \mathbf{u} evaluated at $\mathbf{v} = \mathbf{0}$:

$$\frac{\partial \ell_{adv}(\mathbf{x} + R(\mathbf{v}, \mathbf{U}))}{\partial v_m} = (\mathbf{u}^{(m)})^T \nabla_{\mathbf{x}} \ell_{adv}(\mathbf{x}) = \frac{\partial \ell_{adv}(\mathbf{x})}{\partial \mathbf{u}^{(m)}}.$$

Suppose we have a perfect estimation of gradients with regards to \mathbf{v} , then the rectified perturbations become:

$$\hat{\mathbf{g}} = R(\hat{\mathbf{v}}, \mathbf{U}) = \sum_{m=1}^M \mathbf{u}^{(m)} \cdot \left(\frac{\partial \ell_{adv}(\mathbf{x})}{\partial \mathbf{u}^{(m)}} \right).$$

Since $\frac{\partial \ell_{adv}(\mathbf{x})}{\partial \mathbf{u}^{(m)}}$ is the directional derivative of $\mathbf{u}^{(m)}$, $\mathbf{u}^{(m)} \cdot \frac{\partial \ell_{adv}(\mathbf{x})}{\partial \mathbf{u}^{(m)}}$ is the projection of adversarial gradient on the direction $\mathbf{u}^{(m)}$, if the model is differentiable. Thus, rectified perturbations $\hat{\mathbf{g}}$ is the vector sum of the gradient projection on each patch direction $\mathbf{u}^{(m)}$. Rectified perturbation $\hat{\mathbf{g}}$ can be regarded as the projection of adversarial gradient \mathbf{g} on subspace $\mathbf{s}_{\mathbf{u}}$, as the projection of a vector on a subspace is the vector sum of the vector's projections on the orthogonal basis vectors of the subspace. Figure 3 illustrates a toy example of the projection, where \mathbf{g} is the adversarial gradient, $\hat{\mathbf{g}}$ is the projection of \mathbf{g} on a subspace $\mathbf{s}_{\mathbf{u}}$ that is defined by the direction vectors of two patches: $\mathbf{u}^{(1)}, \mathbf{u}^{(2)}$.

A beneficial property of the projection is that it is the closest vector to the gradient in subspace having $\hat{\mathbf{g}} = \arg \min_{\boldsymbol{\eta}} (\|\mathbf{g} - \boldsymbol{\eta}\|_2), \forall \boldsymbol{\eta} \in \mathbf{s}_{\mathbf{u}} \subseteq \mathcal{S}$. This enables us to only consider the subspace $\mathbf{s}_{\mathbf{u}}$ to find a good approximate of the adversarial gradient \mathbf{g} . In an extreme case where each input dimension (or pixel) is a patch, the rectified perturbation becomes an exact estimate of the original pixel-wise adversarial gradient. However, this will require more queries to the target model, thus suffers from low efficiency. While in a typical case, patch-wise rectification is equivalent to estimating the projection of the adversarial gradient on a subspace $\mathbf{s}_{\mathbf{u}}$ and only requires a reasonable number of queries. On the other hand, better transferability of the tentative perturbations makes the subspace $\mathbf{s}_{\mathbf{u}}$ closer to the true adversarial gradient, and therefore leads to better approximations.

4 EXPERIMENTS

In this section, we provide a comprehensive evaluation of our proposed V-BAD framework and its variants, for both untargeted and targeted video attacks on three benchmark video datasets, against two state-of-the-art video recognition models. We also investigate different choices of tentative perturbations, partitioning methods and estimation methods in an ablation study.

4.1 Experimental Setting

Attack Setting. We consider three benchmark datasets for video recognition: UCF-101 [36], HMDB-51 [21], and Kinetics-400 [18]. For each dataset, we randomly select one test video, from each category, that is correctly classified by the target model. For targeted attack, we randomly choose the target class for each video. For all datasets, we set the perturbation bound to $\epsilon_{adv} = 0.05$ per frame. The query limit, *i.e.*, the maximum number of queries to the target model is set to $Q = 3 \times 10^5$, which is similar to the number of queries required for most black-box image attacks to succeed. We run the attack until an adversarial example is found (attack succeeds) or we reach the query limit. We evaluate different attacking strategies in terms of 1) success rate (SR), the ratio of successful generation of adversarial examples under the perturbation bound within the limited number of queries; and 2) average number of queries (ANQ), required for a successful attack (excluding failed attacks).

Video Recognition Models. We consider two state-of-the-art video recognition models I3D and CNN+LSTM, as our target models to attack. I3D is an inflated 3D convolutional network. We use a Kinetics-400 pretrained I3D and finetune it on other two datasets. We sample frames at 25 frames per second. For CNN+LSTM, we

Table 1: Test Accuracy (%) of the video models.

Model	UCF-101	HMDB-51	Kinetics-400
I3D	91.30	63.73	62.30
CNN+LSTM	76.29	44.38	53.20

Table 2: Results for V-BAD with different: 1) tentative perturbations, 2) partitioning methods, and 3) estimation methods. Symbol “+” indicates fixed methods for untested components of V-BAD. Best results are highlighted in bold.

Components of V-BAD	Method	UCF-101	
		ANQ	SR (%)
Tentative Perturbations + partition: Uniform + estimation: NES	Static	51786	70
	Random	107499	95
	Single	57361	100
Partitioning Method + tentative: Ensemble + estimation: NES	Ensemble	49797	100
	Random	77881	100
	Uniform	49797	100
Estimation Method + tentative: Ensemble + partition: Uniform	FD	61585	70
	NES	49797	100

use an ImageNet pretrained ResNet101 as the CNN feature extractor, then finetune an LSTM on top of it. Input video frames are subsampled by keeping one out of every 5 for CNN+LSTM Model. Note we only consider the RGB part for both video models. At test time, we use 32-frame snippets for UCF-101 and HMDB-51, and 64-frame snippets for Kinetics-400, as it has longer videos. The test accuracy of the two models can be found in Table 1. The accuracy gap between ours and the one reported in [4] is mainly caused by the availability of fewer input frames at test time.

Image Models. We use ImageNet [7] pretrained deep networks as our image model to extract tentative perturbations. Given the difference between our video datasets and the ImageNet, rather than using one model, we use an ensemble of ImageNet models: ResNet50 [10], DenseNet121 [11], and DenseNet169 [11].

V-BAD Setting. For tentative perturbations, we average the perturbations extracted from the three image models to obtain the final perturbations. For partitioning, we use uniform partitioning to get 8×8 patches per frame and set the NES population size of each estimation as 48, which works consistently well across different datasets in terms of both success rate and number of queries. For search variance σ in NES, we set it to 10^{-6} for the targeted attack setting and 10^{-3} for the untargeted attack setting. This is because a targeted attack needs to keep the target class in the top-1 class list to get the score of the targeted class, while the aim of an untargeted attack is to remove the current class from top-1 class position, which allows a large search step.

4.2 Ablation Study of V-BAD

In this section, we evaluate variants of V-BAD, with different types of 1) tentative perturbations, 2) partitioning methods and 3) estimation methods. Experiments were conducted on a subset of 20 randomly selected categories from UCF-101 dataset.

Table 3: Targeted attacks on UCF-101/HMDB-51/Kinetics-400 against I3D/CNN+LSTM models. Best results are in bold.

Target Model	Attack	UCF-101		HMDB-51		Kinetics-400	
		ANQ	SR (%)	ANQ	SR (%)	ANQ	SR (%)
I3D	SR-BAD	67909	96.0	40824	96.1	76986	96.0
	P-BAD	104986	96.0	62744	96.8	98402	96.0
	V-BAD	60687	98.0	34260	96.8	66379	98.0
CNN + LSTM	SR-BAD	147322	45.5	67037	82.4	109314	73.0
	P-BAD	159723	60.4	72697	90.2	117368	85.0
	V-BAD	84294	93.1	44944	98.0	70897	97.0

Tentative Perturbations. We first evaluate V-BAD with the three different types of tentative perturbations: 1) Random, 2) Static, and 3) Transferred. For transferred, we test two different strategies, with either a single image model ResNet-50 (denoted as “Single”) or an ensemble of the three considered image models (denoted as “Ensemble”). The partitioning and estimation methods were set to Uniform and NES respectively. The results can be found in Table 2. A clear improvement can be observed for the use of transferred tentative perturbations compared with static or random perturbations. The number of queries required for successful attacks was dramatically reduced from more than 10^5 to less than 6×10^4 by using only a single image model. This confirms that an image model alone can provide effective guidance for attacking video models. The number was further reduced to around 5×10^4 via the use of an ensemble of three image models. Compared with random perturbations, static perturbations require fewer queries to succeed, with 50% less queries due to the reduced exploration space. However, static perturbations have a much lower success rate (70%) than random perturbations (95%). This indicates that fixed directions can generate adversarial examples faster with restricted exploration space, however, such restrictions may cause the attack easily stuck in some local optima, without the exploration over other directions that could lead to potentially more powerful attacks.

Partitioning Methods. We investigate the two types of partitioning methods introduced in Section 3.4, namely, Random and Uniform. For tentative perturbations, we use the best method found in the previous experiments - Ensemble. Results are reported in Table 2. As can be seen, both partitioning methods can find successful attacks, but the use of uniform partitioning significantly reduces the number of queries by 36%, to around 5×10^4 from 8×10^4 of random partitioning. This is because a tentative perturbations generated from image models often contains certain local patterns, but random partitioning tends to destroy such locality. Recall that partitioning is applied in every step of perturbations, and as such, uniform partitioning can help to maintain stable and consistent patches across different iteration steps. This allows the rectification to make continuous corrections to the same local patches.

Estimation Methods. As we mentioned in Section 3.4, any derivative-free (or black-box) optimization methods can be used to estimate the rectification weights. Here, we compare two of the methods that have been used for black-box image attacks: FD (Finite Difference) [2] and NES (Natural Evolution Strategies) [12]. For a fair comparison, we made some adjustments to the number of patches used by the FD estimator, so that FD and NES require a similar number of queries per update. The results are reported

Table 4: Untargeted attacks on UCF-101/HMDB-51/Kinetics-400 against I3D/CNN+LSTM models. Best results are in bold.

Target Model	Attack	UCF-101		HMDB-51		Kinetics-400	
		ANQ	SR (%)	ANQ	SR (%)	ANQ	SR (%)
I3D	SR-BAD	5143	98.0	1863	100.0	1496	100.0
	P-BAD	11571	98.0	4162	100.0	3167	100.0
	V-BAD	3642	100.0	1152	100.0	1012	100.0
CNN + LSTM	SR-BAD	8674	100.0	684	100.0	1181	100.0
	P-BAD	12628	100.0	1013	100.0	1480	100.0
	V-BAD	784	100.0	197	100.0	293	100.0

in Table 2. NES demonstrates a clear advantage over FD: FD only achieves 70% success rate within the query limit (e.g. 3×10^5), while NES has a 100% success rate. And the number of queries required by successful NES attacks is roughly 20% less than that of FD attacks.

Based on the ablation results, in the following experiments, we set the V-BAD to: ensemble tentative perturbations, uniform partitioning and NES rectification estimator.

4.3 Comparison to Existing Attacks

We compare our V-BAD framework with two existing state-of-the-art black-box image attack methods [2, 12]. Instead of directly applying the two image attacks to videos, we incorporate their logic into our V-BAD framework to obtain two variants of V-BAD.

Baselines. The first baseline method is the pixel-wise adversarial gradient estimation using NES, proposed in [12]. This can be easily achieved by using static tentative perturbations and setting the patches in V-BAD to pixels, *i.e.*, each pixel is a patch. We denote this baseline variant of V-BAD by P-BAD. The NES population size for P-BAD is set to 96, since there are more parameters to estimate. The second baseline method is grouping-based adversarial gradient estimation using FD, proposed by [2]. This method explores random partitioning to reduce the large number of queries required by FD. Accordingly, we use the variant of V-BAD that utilizes static tentative perturbations and random partitioning, and denote it by SR-BAD. Different from its original setting with FD, here we use NES for SR-BAD which was found more efficient in our previous experiments. It is worth mentioning that a dimensionality reduction technique (*e.g.* PCA) was also explored in [2] to project inputs to low dimensional features. Although it increased attack success rate on simple black-white images (*e.g.* MNIST [23]), it did not help attacks on natural images (*e.g.* CIFAR10 [19]), due to the information loss caused by the projection.

Targeted Black-box Video Attack. Comparison results for targeted attacks are reported in Table 3. Among the three methods, V-BAD achieves the best success rates, consistently using least number of queries across the three datasets and two recognition models. Specifically, V-BAD only takes $(3.4 \sim 8.4) \times 10^4$ queries to achieve a success rate of above 93%. Note that this is comparable to state-of-the-art black-box image attacks [12]. Comparing P-BAD and V-BAD, pixel-wise estimation by P-BAD does not seem to yield more effective attacks, whereas on the contrary, partition-based rectifications by V-BAD not only reduces $\sim 50\%$ queries, but also leads to more successful attacks. Compare the performance on different target models, an obvious degradation of performance can be observed on CNN+SLTM model. This is because CNN+SLTM has

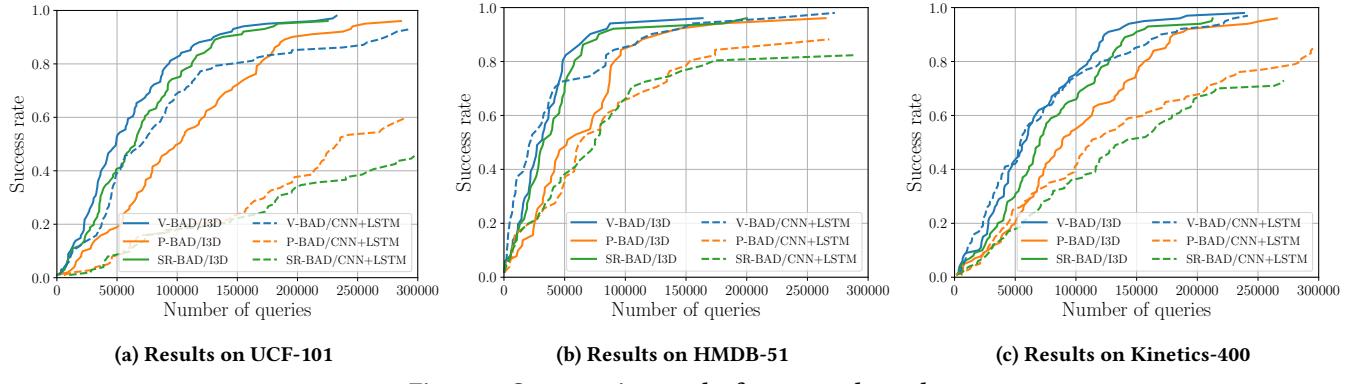


Figure 4: Comparative results for targeted attack.

Table 5: Cosine similarity between varies tentative or rectified gradients and the actual gradient.

Tentative	Static	Random	Transferred
Cosine	7.177×10^{-5}	-1.821×10^{-5}	-2.743×10^{-4}
Rectified	SR-BAD	P-BAD	V-BAD
Cosine	3.480×10^{-3}	3.029×10^{-3}	4.661×10^{-3}

a lower accuracy than I3D, making it relatively robust to targeted attacks (not for untargeted attacks), an observation that is consistent with findings in [38]. However, this impact is the least significant on V-BAD where the accuracy decreases less than 5%, while the accuracy of P-BAD and SR-BAD has a huge drop, especially on UCF-101 (from 96.0% to 45.5%). This is further illustrated in Figure 4, showing change in success rate with the number of queries. This can probably be explained by the better transferability of transferred tentative perturbations on CNN+LSTM than I3D due to the similar 2D CNN used in CNN+LSTM video model. As in Figure 4, the advantage of better transferability even overcomes the low accuracy drawback of CNN+LSTM on HMDB-51 and Kinetics-400 dataset: V-BAD/CNN+LSTM is above V-BAD/I3D for the first 4×10^4 queries. A targeted video adversarial examples generated by V-BAD is illustrated in Figure 1, where video on the top is the original video with the correct class and video at the bottom is the video adversarial example misclassified as the adversarial class.

Untargeted Black-box Video Attack. Results for untargeted attacks are reported in Table 4. Compared with targeted attacks, untargeted attacks are much easier to achieve, requiring only $\sim 10\%$ queries of targeted attacks. Compared with other baselines, V-BAD is the most effective and efficient attack across all datasets and recognition models. It only takes a few hundred queries for V-BAD to completely break the CNN+LSTM model.

Both attacks indicate that video models are as vulnerable as image models to black-box adversarial attacks. This has serious implications for the video recognition community to consider.

4.4 Gradient Estimate Quality

We further explore the quality of various tentative perturbations and rectified perturbations generated by different variants of V-BAD. We measure the perturbations quality by calculating the cosine similarity between the ground-truth adversarial gradient and the tentative/rectified perturbations. The results are based on

20 random runs of the attacks on 50 videos randomly chosen from UCF-101, and are reported in Table 5. Consistent with the comparison experiments, V-BAD generates the best gradient estimates and P-BAD has the worst estimation quality. All the rectified perturbations are much better than the tentative perturbations. This verifies that tentative perturbations can be significantly improved by proper rectification. One interesting observation is that the transferred tentative perturbations (from an ensemble of ImageNet models) have a large negative cosine similarity, which is opposite to our design. One explanation could be that there is a huge gap between the image model and the video model. However, note that while the transferred perturbations is opposite to the gradient, it serves as a good initialization and yields better gradient estimation after rectification. It is noteworthy that there is still a considerable gap between the gradient estimate and the actual gradient. From one angle, it reflects that we do not need very accurate gradient estimation to generate adversarial examples. From another angle, it suggests that black-box attack based on gradient estimation has great scope for further improvement.

5 CONCLUSION

In this paper, we investigated the problem of black-box adversarial attack against video recognition models, and proposed the first framework, V-BAD, for the generation of video adversarial examples through only queries to a video model. To address efficiency issues caused by the high dimensionality of videos, we decoupled adversarial gradient estimation into a two-step process: tentative perturbations transfer followed by patch-wise rectification with NES estimation method. We demonstrated the effectiveness and efficiency of V-BAD by attacking two state-of-the-art video recognition models on three benchmark video datasets. Compared with existing black-box methods, V-BAD achieved high success rate using significantly less queries for both targeted and untargeted attacks. Our results suggest that video models are also highly vulnerable to black-box adversarial attacks, and that effective defenses for video models will need to be developed for secure video recognition.

ACKNOWLEDGMENTS

This work was supported in part by National Key Research and Development Program of China under Grant 2018YFB1004300 and National Natural Science Foundation of China under Grant 61622204.

REFERENCES

- [1] Anish Athalye, Nicholas Carlini, and David Wagner. 2018. Obfuscated gradients give a false sense of security: Circumventing defenses to adversarial examples. In *ICML*.
- [2] Arjun Nitin Bhagoji, Warren He, Bo Li, and Dawn Song. 2018. Practical black-box attacks on deep neural networks using efficient query mechanisms. In *ECCV*.
- [3] Nicholas Carlini and David Wagner. 2017. Towards evaluating the robustness of neural networks. In *S&P*.
- [4] Joao Carreira and Andrew Zisserman. 2017. Quo vadis, action recognition? a new model and the kinetics dataset. In *CVPR*.
- [5] Pin-Yu Chen, Yash Sharma, Huan Zhang, Jinfeng Yi, and Cho-Jui Hsieh. 2018. EAD: elastic-net attacks to deep neural networks via adversarial examples. In *AAAI*.
- [6] Pin-Yu Chen, Huan Zhang, Yash Sharma, Jinfeng Yi, and Cho-Jui Hsieh. 2017. Zoo: Zeroth order optimization based black-box attacks to deep neural networks without training substitute models. In *AISec*.
- [7] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. 2009. Imagenet: A large-scale hierarchical image database. In *CVPR*.
- [8] Kevin Eykholt, Ivan Evtimov, Earlene Fernandes, Bo Li, Amir Rahmati, Chaowei Xiao, Atul Prakash, Tadayoshi Kohno, and Dawn Song. 2018. Robust Physical-World Attacks on Deep Learning Visual Classification. In *CVPR*.
- [9] Ian J. Goodfellow, Jonathon Shlens, and Christian Szegedy. 2015. Explaining and harnessing adversarial examples. In *ICLR*.
- [10] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2016. Deep residual learning for image recognition. In *CVPR*.
- [11] Gao Huang, Zhuang Liu, Laurens Van Der Maaten, and Kilian Q Weinberger. 2017. Densely connected convolutional networks. In *CVPR*.
- [12] Andrew Ilyas, Logan Engstrom, Anish Athalye, and Jessy Lin. 2018. Black-box adversarial attacks with limited queries and information. *ICML*.
- [13] Shuiwang Ji, Wei Xu, Ming Yang, and Kai Yu. 2013. 3D convolutional neural networks for human action recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 35, 1 (2013), 221–231.
- [14] Yu-Gang Jiang, Minjun Li, Xi Wang, Wei Liu, and Xian-Sheng Hua. 2018. Deep-Product: Mobile Product Search With Portable Deep Features. *ACM Transactions on Multimedia Computing, Communications and Applications (TOMM)* 14, 2 (2018), 50:1–50:18.
- [15] Yu-Gang Jiang, Zuxuan Wu, Jinhui Tang, Zechao Li, Xiangyang Xue, and Shih-Fu Chang. 2018. Modeling Multimodal Clues in a Hybrid Deep Learning Framework for Video Classification. *IEEE Transactions on Multimedia (TMM)* 20, 11 (2018), 3137–3147.
- [16] Yu-Gang Jiang, Zuxuan Wu, Jun Wang, Xiangyang Xue, and Shih-Fu Chang. 2018. Exploiting feature and class relationships in video categorization with regularized deep neural networks. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 40, 2 (2018), 352–364.
- [17] Andrej Karpathy, George Toderici, Sanketh Shetty, Thomas Leung, Rahul Sukthankar, and Li Fei-Fei. 2014. Large-scale video classification with convolutional neural networks. In *CVPR*.
- [18] Will Kay, Joao Carreira, Karen Simonyan, Brian Zhang, Chloe Hillier, Sudheendra Vijayanarasimhan, Fabio Viola, Tim Green, Trevor Back, Paul Natsev, et al. 2017. The kinetics human action video dataset. *arXiv preprint arXiv:1705.06950* (2017).
- [19] Alex Krizhevsky and Geoffrey Hinton. 2009. *Learning multiple layers of features from tiny images*. Technical Report. Citeseer.
- [20] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton. 2012. Imagenet classification with deep convolutional neural networks. In *NIPS*.
- [21] Hildegard Kuehne, Hueihan Jhuang, Estibaliz Garrote, Tomaso Poggio, and Thomas Serre. 2011. HMDB: a large video database for human motion recognition. In *ICCV*.
- [22] Alexey Kurakin, Ian Goodfellow, and Samy Bengio. 2016. Adversarial examples in the physical world. *arXiv preprint arXiv:1607.02533* (2016).
- [23] Yann LeCun, Bernhard E Boser, John S Denker, Donnie Henderson, Richard E Howard, Wayne E Hubbard, and Lawrence D Jackel. 1990. Handwritten digit recognition with a back-propagation network. In *NIPS*.
- [24] Shasha Li, Ajaya Neupane, Sujoy Paul, Chengyu Song, Srikanth V Krishnamurthy, Amit K Roy Chowdhury, and Ananthram Swami. 2018. Adversarial perturbations against real-time video classification systems. *arXiv preprint arXiv:1807.00458* (2018).
- [25] Sheng Liu, Zhou Ren, and Junsong Yuan. 2018. SibNet: Sibling Convolutional Encoder for Video Captioning. In *ACM MM*.
- [26] Xingjun Ma, Bo Li, Yisen Wang, Sarah M. Erfani, Sudanthi Wijewickrema, Michael E. Houle, Grant Schoenebeck, Dawn Song, and James Bailey. 2018. Characterizing adversarial subspaces using local intrinsic dimensionality. In *ICLR*.
- [27] Aleksander Madry, Aleksandar Makelov, Ludwig Schmidt, Dimitris Tsipras, and Adrian Vladu. 2018. Towards deep learning models resistant to adversarial attacks. In *ICLR*.
- [28] Seyed-Mohsen Moosavi-Dezfooli, Alhussein Fawzi, Omar Fawzi, and Pascal Frossard. 2017. Universal adversarial perturbations. In *CVPR*.
- [29] Seyed-Mohsen Moosavi-Dezfooli, Alhussein Fawzi, and Pascal Frossard. 2016. Deepfool: a simple and accurate method to fool deep neural networks. In *CVPR*.
- [30] Anh Nguyen, Jason Yosinski, and Jeff Clune. 2015. Deep neural networks are easily fooled: High confidence predictions for unrecognizable images. In *CVPR*.
- [31] Nicolas Papernot, Patrick McDaniel, and Ian Goodfellow. 2016. Transferability in machine learning: from phenomena to black-box attacks using adversarial samples. *arXiv preprint arXiv:1605.07277* (2016).
- [32] Nicolas Papernot, Patrick McDaniel, Ian Goodfellow, Somesh Jha, Z Berkay Celik, and Ananthram Swami. 2017. Practical black-box attacks against machine learning. In *ASIACCS*.
- [33] Nicolas Papernot, Patrick McDaniel, Somesh Jha, Matt Fredrikson, Z Berkay Celik, and Ananthram Swami. 2016. The limitations of deep learning in adversarial settings. In *EuroS&P*.
- [34] Roberto Rey-de Castro and Herschel Rabitz. 2018. Targeted nonlinear adversarial perturbations in images and videos. *arXiv preprint arXiv:1809.00958* (2018).
- [35] Tim Salimans, Jonathan Ho, Xi Chen, Szymon Sidor, and Ilya Sutskever. 2017. Evolution strategies as a scalable alternative to reinforcement learning. *arXiv preprint arXiv:1703.03864* (2017).
- [36] Khurram Soomro, Amir Roshan Zamir, and Mubarak Shah. 2012. UCF101: A dataset of 101 human actions classes from videos in the wild. *arXiv preprint arXiv:1212.0402* (2012).
- [37] Christian Szegedy, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian Goodfellow, and Rob Fergus. 2014. Intriguing properties of neural networks. In *ICLR*.
- [38] Dimitris Tsipras, Shibani Santurkar, Logan Engstrom, Alexander Turner, and Aleksander Madry. 2019. Robustness may be at odds with accuracy. In *ICLR*.
- [39] Yisen Wang, Xingjun Ma, James Bailey, Jinfeng Yi, Bowen Zhou, and Quanquan Gu. 2019. On the Convergence and Robustness of Adversarial Training. In *ICML*. 6586–6595.
- [40] Xingxing Wei, Jun Zhu, and Hang Su. 2018. Sparse adversarial perturbations for videos. *arXiv preprint arXiv:1803.02536* (2018).
- [41] Daan Wierstra, Tom Schaul, Tobias Glasmachers, Yi Sun, Jan Peters, and Jürgen Schmidhuber. 2014. Natural evolution strategies. *Journal of Machine Learning Research* 15, 1 (2014), 949–980.
- [42] Zuxuan Wu, Yu-Gang Jiang, Xi Wang, Hao Ye, and Xiangyang Xue. 2016. Multi-stream multi-class fusion of deep networks for video classification. In *ACM MM*.
- [43] Ziwei Yang, Yahong Han, and Zheng Wang. 2017. Catching the temporal regions-of-interest for video captioning. In *ACM MM*.
- [44] Joe Yue-Hei Ng, Matthew Hausknecht, Sudheendra Vijayanarasimhan, Oriol Vinyals, Rajat Monga, and George Toderici. 2015. Beyond short snippets: Deep networks for video classification. In *CVPR*.
- [45] Rui-Wei Zhao, Zuxuan Wu, Jianguo Li, and Yu-Gang Jiang. 2017. Learning semantic feature map for visual content recognition. In *ACM MM*.